

The Database as the Trusted Computing Base for LLM Agents: Experience from TGMS

Xiaofei Zhang

University of Memphis

Memphis, TN, USA

xiaofei.zhang@memphis.edu

ABSTRACT

Placeholder: existing agent–database interfaces ask an unreliable model to generate a query, interpret partial results, perform additional computation, and state conclusions. We argue agent-native databases should instead expose a small, typed, costed operator algebra and return evidence-carrying results whose downstream claims can be mechanically verified.

- 1 THE PROBLEM AND THE POSITION
- 2 AGENT-NATIVE DATABASE
ARCHITECTURE
- 3 EVIDENCE AND BELIEF SEMANTICS
- 4 THE TGMS PROTOTYPE AND LESSONS
- 5 PRELIMINARY EVIDENCE
- 6 RESEARCH AGENDA