

TGMS: An Agent-Native Bi-Temporal Graph Management System

Verified Temporal Operators as Tools, and Trace-Grounded Answer Verification

Xiaofei Zhang
University of Memphis
xiaofei.zhang@memphis.edu

July 2026 · system description preprint (v1) · <https://github.com/zxf-work/tgms>*

Abstract

Large language model (LLM) agents are unreliable at exactly the skills that temporal graph analytics requires: arithmetic, identifier fidelity, and asserting only what evidence supports. We present TGMS, a bi-temporal property graph management system whose query surface is not a query language but a small algebra of thirteen *verified temporal operators* exposed to agents as tools—typed, deterministic, bounded, cost-guarded, and bi-temporal by default—paired with a Planner–Executor–Verifier agent layer in which the LLM only plans and reports, while every numeric, entity, ordering, and pattern claim in a final answer is machine-checked against the content-addressed execution trace that produced it. The bi-temporal substrate (valid time $\times$ transaction time) distinguishes *evolution* from *correction*, enabling belief-state queries (“as of transaction time T , what did we believe?”) that no snapshot or retrieval baseline can express. On a development benchmark over a real communication network, with open-source models served on a single 24 GB GPU, the operator-backed agent reaches 0.41 pooled exact match versus 0.05–0.18 for vector-RAG, static-graph-RAG, and text-to-Cypher baselines, and 0.67 on correction probes where all baselines score zero; the claim verifier detects 500/500 injected false claims at zero false positives. We report two design findings from live small-model traffic: operator *output contracts* (statically rejecting invented result paths) converted execution success from 0.00 to 1.00 on probe tasks, and verification must track *evidence completeness*—a model can compute correct arithmetic over a truncated result page. Code, benchmark, and an auditable trace viewer are open source under Apache-2.0.

1 Introduction

Temporal graphs—communication and transaction logs, evolving knowledge bases—pose two problems for LLM agents that static text corpora do not. First, **time is compositional**: “among accounts reachable from X in February, how many cyclic triangles closed within a day?” chains time-respecting reachability [25] into δ -motif counting [17]. Serializing edges into text and retrieving chunks destroys precisely the structure the question quantifies over. Second, **belief changes**: real stores are corrected after the fact, and answering “what did we believe on March 1” requires separate valid-time and transaction-time axes [24] and the discipline to distinguish evolution (the edge ended) from correction (we were wrong). Snapshot representations cannot express the question.

LLMs are additionally unreliable at arithmetic, prone to inventing identifiers, and prone to asserting numbers absent from evidence. TGMS’s position is architectural: give the model *no*

*This preprint accompanies TGMS v0.1.0 and reports development-split results; a revised version will accompany the pre-registered frozen-test campaign.




Figure 1: TGMS architecture. The LLM only plans and reports; operators compute; the verifier checks every claim against the execution trace.

opportunity to do any of those things. Identifiers must come from a resolution operator or the task input (enforced statically); arithmetic must go through a **compute** operator; every asserted number must trace to a content-addressed result digest, and a verifier checks that it does.

This preprint describes the system design and implementation (Phase 1–2 of a longer research program), the evaluation methodology, first development-split measurements with open-source models, and two design findings that emerged from live model traffic. We pre-registered acceptance thresholds before measurement and report against them.

2 Bi-temporal substrate

Every logical node and edge is an identity plus a set of *versions* carrying half-open valid-time $[vt_s, vt_e)$ and transaction-time $[tt_s, tt_e)$ intervals (int64 epoch microseconds; open end = 2^{62}). The write API is three bi-temporal verbs—**assert** (new belief, with interval carving of overlapping prior belief), **retract** (evolution), **correct** (correction)—plus a bulk ingest path for instantaneous event streams. All writes pass through an append-only write-ahead event log, the single source of truth: replaying it into any backend reproduces a byte-identical store digest (tested across Kùzu [6] and DuckDB [19] adapters). Update semantics are implemented once, above a small adapter interface, so backends cannot diverge; a hybrid logical clock guarantees strictly monotonic transaction times; failed batches remain in the log, where replay deterministically re-fails and skips them. Version rows reserve **source/provenance_ref** columns so future agent write-back requires no migration.

Property-based tests over random assert/retract/correct interleavings check the core invariants: believed versions of one identity have pairwise-disjoint valid intervals *at every historical transaction time*, and—the signature property—**bi-temporal immutability**: any operator result pinned to a past **as_of_tt** is byte-identical before and after later corrections. Enforcing the latter required two non-obvious rules: result digests must exclude current-belief metadata, and returned versions must censor belief-closure timestamps later than the pinned transaction time.

Table 1: The operator algebra. Every operator additionally takes `as_of_tt`, `limit/cursor`, and declares its output fields.

	Operator	Semantics
O1	<code>entity_history</code>	version list of a node under a belief state
O2	<code>snapshot_subgraph</code>	k -hop (≤ 3) neighborhood of the snapshot $G(t, tt)$
O3	<code>diff_snapshots</code>	adds/removes/changes between $G(t_1)$ and $G(t_2)$, from version intervals
O4	<code>temporal_reachability</code>	earliest arrival over time-respecting paths; exact under wait caps
O5	<code>temporal_paths</code>	up to $k \leq 20$ node-simple time-respecting paths
O6/O7	<code>count/find_temporal_motifs</code>	δ -temporal motif catalog [17], exact, order (t, eid)
O8	<code>graph_metric_timeseries</code>	bucketed node/edge/degree/reciprocity series
O9	<code>burst_detection</code>	rolling z-score / trailing-median outlier buckets
O10	<code>neighborhood_evolution</code>	neighbors gained/lost + degree series
O11	<code>co_active</code>	Allen-relation interval join between edge selections
O12	<code>resolve_entities</code>	name/uid lookup — the only legal identifier source
O13	<code>compute</code>	count/sum/min/max/top- k /filter/interval relation over step outputs

3 A verified operator algebra

The query surface is the thirteen operators of Table 1; `compute` keeps arithmetic out of the LLM entirely, and `resolve_entities` is the only legal source of identifiers.

Five design rules bind every operator. **Typed:** inputs and outputs validate against JSON Schemas generated from one registry, which also emits the tool definitions. **Deterministic:** same store and arguments imply byte-identical canonical output carrying a SHA-256 `result_digest`; floats are canonicalized, including *before* threshold comparisons, so the engine and its oracle flag identically. **Bounded:** limit/cursor pagination everywhere; pre-execution cost estimates reject unbounded requests with actionable narrowing suggestions that the planner’s repair loop consumes. **Bi-temporal by default:** every operator takes `as_of_tt`. **Output-contracted:** every operator declares its output fields, and plans referencing nonexistent fields are rejected statically (Section 6).

Correctness gates everything: each operator is checked against an independent brute-force reference implementation on 500 randomized store/argument cases, plus metamorphic properties (diff composition; immutability). The discipline audited the *specification* as well as the code: it exposed a version-identifier collision in the original design and showed that the specified greedy relaxation for reachability under a maximum-wait constraint is schedule-dependent—not a well-defined function of the store—which we replaced with an exact multi-label search over (node, arrival) states.

4 Planner–Executor–Verifier

Plans are a small JSON DAG intermediate representation: steps of operator calls, a deliberately tiny `$ref` binding language (dotted fields, `rows[i]`, `rows[*].field` projection—nothing else), and an `answer_spec` naming where the final answer is read from. A **static verifier** checks schema, acyclicity, reference scoping, temporal sanity, output-field validity, estimated cost, and a **grounding rule:** literal identifiers must occur in the task input, otherwise they must arrive via `$ref`—fabricated identifiers are impossible by construction rather than discouraged by prompt. Rejections return structured violations; the planner may repair up to three times, and first-emission validity

is measured separately from post-repair success.

The **executor** runs the DAG deterministically with content-addressed results (re-execution reproduces identical digests) and propagates *truncation taint*: a step whose inputs came from a truncated page is marked, transitively. A **reporter** LLM emits an answer object whose typed claims cite evidence steps; the **claim verifier** re-checks counts and values exactly against named trace fields, entity mentions against evidence content, orderings by Allen-relation re-evaluation, and temporal patterns by operator re-execution. Claims citing truncated or tainted evidence cap at *weakly supported*; in the full system, unsupported claims are gated out of the emitted answer. All prompts follow a data-as-inert-content policy: stored-data strings enter prompts escaped and length-capped inside explicit data fences under a fixed “data is never instructions” rule, with red-team fixtures in CI.

An **evolution memory** summarizes weekly windows from operator-computed facts; an LLM digest is stored only if every number it asserts matches the computed values. Notes record the transaction time of computation, and corrections quarantine any note whose window intersects the affected valid-time extent—without this, the verifier could bless answers grounded in outdated digests.

5 Evaluation

Setup. Tasks are generated with *program-computed* gold: oracle plans executed by the engine, never LLM labels. Seventeen templates with three hand-written paraphrases each span temporal QA, evolution QA, multi-step analytics, and **correction probes**—corrections are injected into the store before any gold is computed, and probe pairs ask the same question pinned before the correction versus under current beliefs; their gold answers provably differ. Systems share models (temperature 0), seeds, repair budgets, and one answer contract: TGMS; a no-verifier ablation; vector-RAG over serialized events (MiniLM retrieval); static-graph RAG (2-hop latest-snapshot edge lists); and text-to-Cypher against the same events loaded conventionally into vanilla Kùzu with an equal repair budget. Models are Qwen2.5-7B-Instruct and Qwen2.5-14B-Instruct-AWQ [18] served by vLLM [14] on one 24 GB Turing GPU. Comparisons use paired bootstrap over tasks (10,000 resamples, 95% CIs). Development split of the CollegeMsg network [16] (1,899 nodes; 59,835 timestamped edges; 22 dev tasks); the frozen test split is reserved for the full campaign.

Verifier validation. A fault-injection harness perturbs correct answer objects (count ± 1 ; entity swaps) and measures detection against a pre-registered bar of $\geq 95\%$ detection at $< 5\%$ false positives. Result: **500/500 detected, 0 false positives**; emitted answers carry an unsupported-claim rate of 0.000.

End-to-end results. Table 2 summarizes the development-split campaign.

Three observations. (1) The **bi-temporal advantage is model-size-independent**: correction probes score 0.67 at both 7B and 14B while every baseline scores zero at 14B—the belief-time dimension is not representable in their inputs. (2) **Scaling is steep at small sizes**: 7B→14B triples pooled exact match and lifts first-emission plan validity from 0.08–0.17 to 0.33–0.75 by family; the repair loop is worth roughly a tripling of execution success at 7B. (3) Operators meet all latency targets at 10^6 events on a 40-core server (snapshot subgraph 98 ms; global diff 163 ms; reachability 63–244 ms; series ~ 155 ms, p50), after profiling showed that materializing unused string columns dominated scans.

Table 2: Pooled exact match on the CollegeMsg dev split (22 tasks; single seed; temperature 0). Paired-bootstrap 95% CIs for TGMS—baseline: vs. B2 [+0.18, +0.59] and vs. B1 [+0.09, +0.55] (both significant); vs. B5 [0.00, +0.46] (borderline at $n=22$).

	TGMS	B1 vector-RAG	B2 static-RAG	B5 text-to-Cypher
Qwen2.5-7B, all families	0.136	0.045	0.045	0.091
Qwen2.5-14B, all families	0.409	0.091	0.045	0.182
Qwen2.5-14B, correction probes	0.67	0.00	0.00	0.00

6 Findings from live model traffic

The operator manual must be a contract. Our first live matrix run produced zero execution success across all systems-under-test rows: models emit plans that are statically valid yet reference nonexistent *output* paths (e.g., `s2.count` on an operator that emits `rows_total`). No input-schema rigor covers this. Declaring each operator’s output fields in the registry, rejecting bad paths statically, and returning the real field list in the repair payload moved probe-task execution success from 0.00 to 1.00 in one change.

Verify evidence completeness, not just values. A 14B model planned reachability without raising the page limit, computed a count over the truncated 100-row page (true answer: 343), and the claim verified as *supported*—correct arithmetic over incomplete evidence. Truncation now taints dependent steps transitively, and claims citing tainted evidence cap at weakly supported. We conjecture a family of such “evidence-quality” dimensions (cost-narrowed windows, sampling) that verification systems should track as a lattice rather than a boolean.

Serving envelopes are a fairness dimension. The standard vector-RAG configuration (top- $k=20$ chunks of 256 serialized events) assumes frontier context windows: numeric-dense event text tokenizes at ~ 0.65 tokens/character, so even $k=3$ overflows a 28k window once the decoder’s output reservation is charged. Its feasible best on this hardware ($k=1$) covers 0.4% of the corpus per question. Baseline fairness must be defined per serving envelope; the operator system’s per-task consumption ($\sim 8\text{--}10\text{k}$ tokens) is corpus-size-independent.

7 Related work

Temporal and bi-temporal graph databases. Bi-temporal semantics — valid versus transaction time, evolution versus correction — are classical [24]. Recent systems bring native temporal support to property graphs: AeonG [11] adds built-in temporality with a current/historical storage split (mirrored in our substrate), Gradoop’s TPGM supports distributed temporal graph analytics [21], and T-GQL proposes a temporal query language [3]. These systems target human-written queries in a general language. TGMS deliberately does not: its surface is a small, closed algebra whose contracts (types, bounds, cost ceilings, output fields) exist so that *language models* can plan over it and *verifiers* can check the results—and, to our knowledge, none of these systems expose transaction-time reasoning to agents or verify answers against execution.

Graph-augmented and agent memory for LLMs. GraphRAG summarizes corpus-derived graphs for query-focused retrieval [5]; HippoRAG uses graph structure for long-term memory consolidation [9]; Zep/Graphiti is closest to us in adopting a bi-temporal knowledge graph for agent

memory [20]. The shared paradigm is *retrieval into context*: structure is serialized and the model reasons over it in-context. TGMS inverts this: the model never sees raw structure at scale; operators compute, and the model composes and reports. Our correction probes make the difference measurable — retrieval-based systems score zero regardless of model size because their inputs do not contain belief history — and our memory layer adds a guarantee the retrieval paradigm lacks: digests are quarantined when corrections touch their windows, so verification never blesses answers grounded in retracted facts.

Tool use, planning, and program-aided reasoning. Tool-augmented LLMs [23], interleaved reasoning and acting [26], and compiled parallel function-calling plans [13] established that models can orchestrate external calls; program-aided approaches showed that delegating computation to an interpreter beats in-context arithmetic [8, 2]. TGMS applies both lessons with a domain-specific twist: the plan is a constrained DAG IR over a *closed, semantically verified* operator set (each operator oracle-tested on hundreds of randomized cases), and the executable surface itself enforces grounding and output contracts statically. Our finding that output contracts moved small-model execution success from 0.00 to 1.00 suggests tool *result* schemas deserve the same first-class treatment the literature gives argument schemas.

Natural-language interfaces to databases. Text-to-SQL and text-to-Cypher map questions directly to queries [27]; our B5 baseline instantiates this fairly (same models, repair budget, answer contract) on a vanilla property graph. Beyond the accuracy gap on temporal composition, the structural contrast is auditability: generated-query answers can only be re-executed wholesale, whereas TGMS answers carry per-claim evidence digests that a verifier checks mechanically. Temporal KGQA benchmarks [22] probe time-aware questions over knowledge graphs but assume a fixed, correct KG; our correction probes target the orthogonal axis of *belief revision*, which those benchmarks—like the systems above—cannot express.

Faithfulness and claim verification. Post-hoc attribution and revision [7], atomic-fact scoring [15], and self-verification chains [4] check generated text against retrieved or model-internal evidence, probabilistically. TGMS’s verifier is narrower and stronger: claims cite content-addressed results of deterministic operators, so checking is exact (500/500 injected faults caught at zero false positives), and evidence *completeness* is tracked through the plan (truncation taint)—a dimension that text-attribution methods do not model.

Temporal graph analysis. Our operator semantics follow the temporal-networks literature: time-respecting paths [25, 10] and δ -temporal motifs [17], with datasets in the SNAP/TGB tradition [16, 12]. TGMS packages these algorithms behind agent-usable contracts with exactness guarantees and explicit cost refusal instead of silent sampling—and, for reachability under wait constraints, replaces the schedule-dependent greedy relaxation with an exact multi-label semantics.

Positioning. Across these lines of work, TGMS occupies an unclaimed intersection: (i) database-grade temporal semantics including the transaction-time axis, (ii) an agent-facing computation surface with machine-checkable contracts rather than a query language or a retrieval index, and (iii) verification that is exact because the evidence is deterministic. The correction-probe benchmark operationalizes (i), the output-contract finding motivates (ii), and the fault-injection results

validate (iii); the tool server speaks MCP [1], making the combination available to any agent framework.

8 Limitations and roadmap

Reported numbers are development-split, single-seed, single-dataset; the frozen test campaign (more datasets, models, seeds; pre-registered thresholds recorded before any test-split run) is the immediate next step, along with an end-to-end hallucination-reduction readout against the no-verifier ablation, guided/constrained decoding for small-model plan validity, and a concurrency study of many agents sharing one store. Pattern-type claims are reported, not gated. Agent write-back with provenance is designed for (schema and log fields exist) but not enabled.

9 Conclusion

Making a temporal graph store *agent-native*—typed, deterministic, bounded operators; grounding and output contracts; trace-grounded verification; bi-temporal semantics—turns LLM answer quality from a prompting problem into a systems problem with measurable contracts. The early evidence is that the contracts do the heavy lifting: they are what small open models need to be useful, and what makes answers auditable at any model size.

Artifacts. Code, benchmark generator, task suites, the guided demo, and an auditable trace viewer: <https://github.com/zxf-work/tgms> (Apache-2.0).

References

- [1] Anthropic. Model context protocol. <https://modelcontextprotocol.io>, 2024.
- [2] Wenhua Chen, Xueguang Ma, Xinyi Wang, and William W. Cohen. Program of thoughts prompting: Disentangling computation from reasoning for numerical reasoning tasks. *Transactions on Machine Learning Research*, 2023.
- [3] Ariel Debrouvier, Eliseo Parodi, Matías Perazzo, Valeria Soliani, and Alejandro Vaisman. A model and query language for temporal graph databases. *The VLDB Journal*, 30:825–858, 2021.
- [4] Shehzaad Dhuliawala, Mojtaba Komeili, Jing Xu, Roberta Raileanu, Xian Li, Asli Celikyilmaz, and Jason Weston. Chain-of-verification reduces hallucination in large language models, 2023. arXiv:2309.11495.
- [5] Darren Edge, Ha Trinh, Newman Cheng, Joshua Bradley, Alex Chao, Apurva Mody, Steven Truitt, and Jonathan Larson. From local to global: A graph RAG approach to query-focused summarization, 2024. arXiv:2404.16130.
- [6] Xiyang Feng, Guodong Jin, Ziyi Chen, Chang Liu, and Semih Salihoğlu. Kùzu graph database management system. In *Conference on Innovative Data Systems Research (CIDR)*, 2023.
- [7] Luyu Gao, Zhuyun Dai, Panupong Pasupat, Anthony Chen, Arun Tejasvi Chaganty, Yicheng Fan, Vincent Y. Zhao, Ni Lao, Hongrae Lee, Da-Cheng Juan, and Kelvin Guu. RARR:

- Researching and revising what language models say, using language models. In *Proceedings of the 61st Annual Meeting of the Association for Computational Linguistics (ACL)*, 2023.
- [8] Luyu Gao, Aman Madaan, Shuyan Zhou, Uri Alon, Pengfei Liu, Yiming Yang, Jamie Callan, and Graham Neubig. PAL: Program-aided language models. In *International Conference on Machine Learning (ICML)*, 2023.
 - [9] Bernal Jiménez Gutiérrez, Yiheng Shu, Yu Gu, Michihiro Yasunaga, and Yu Su. HippoRAG: Neurobiologically inspired long-term memory for large language models. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2024.
 - [10] Petter Holme and Jari Saramäki. Temporal networks. *Physics Reports*, 519(3):97–125, 2012.
 - [11] Jiamin Hou, Zhanhao Zhang, Zhouyu Wang, Yongjun Zhang, Wei Lu, Anqun Pan, and Xiaoyong Du. Aeong: An efficient built-in temporal support in graph databases. *Proceedings of the VLDB Endowment*, 17(6):1515–1527, 2024.
 - [12] Shenyang Huang, Farimah Poursafaei, Jacob Danovitch, Matthias Fey, Weihua Hu, Emanuele Rossi, Jure Leskovec, Michael Bronstein, Guillaume Rabusseau, and Reihaneh Rabbany. Temporal graph benchmark for machine learning on temporal graphs. In *Advances in Neural Information Processing Systems (NeurIPS), Datasets and Benchmarks Track*, 2023.
 - [13] Sehoon Kim, Suhong Moon, Ryan Tabrizi, Nicholas Lee, Michael W. Mahoney, Kurt Keutzer, and Amir Gholami. An LLM compiler for parallel function calling. In *International Conference on Machine Learning (ICML)*, 2024.
 - [14] Woosuk Kwon, Zhuohan Li, Siyuan Zhuang, Ying Sheng, Lianmin Zheng, Cody Hao Yu, Joseph E. Gonzalez, Hao Zhang, and Ion Stoica. Efficient memory management for large language model serving with pagedattention. In *Proceedings of the 29th Symposium on Operating Systems Principles (SOSP)*, 2023.
 - [15] Sewon Min, Kalpesh Krishna, Xinxu Lyu, Mike Lewis, Wen-tau Yih, Pang Wei Koh, Mohit Iyyer, Luke Zettlemoyer, and Hannaneh Hajishirzi. FActScore: Fine-grained atomic evaluation of factual precision in long form text generation. In *Proceedings of the 2023 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2023.
 - [16] Pietro Panzarasa, Tore Opsahl, and Kathleen M. Carley. Patterns and dynamics of users’ behavior and interaction: Network analysis of an online community. *Journal of the American Society for Information Science and Technology*, 60(5):911–932, 2009.
 - [17] Ashwin Paranjape, Austin R. Benson, and Jure Leskovec. Motifs in temporal networks. In *Proceedings of the Tenth ACM International Conference on Web Search and Data Mining (WSDM)*, pages 601–610, 2017.
 - [18] Qwen Team. Qwen2.5 technical report, 2024. arXiv:2412.15115.
 - [19] Mark Raasveldt and Hannes Mühleisen. Duckdb: An embeddable analytical database. In *Proceedings of the 2019 International Conference on Management of Data (SIGMOD)*, pages 1981–1984, 2019.
 - [20] Preston Rasmussen, Pavlo Paliychuk, Travis Beauvais, Jack Ryan, and Daniel Chalef. Zep: A temporal knowledge graph architecture for agent memory, 2025. arXiv:2501.13956.

- [21] Christopher Rost, Kevin Gomez, Matthias Täschner, Philip Fritzsche, Lucas Schons, Lukas Christ, Timo Adameit, Martin Junghanns, and Erhard Rahm. Distributed temporal graph analytics with GRADOOP. *The VLDB Journal*, 31:375–401, 2022.
- [22] Apoorv Saxena, Soumen Chakrabarti, and Partha Talukdar. Question answering over temporal knowledge graphs. In *Proceedings of the 59th Annual Meeting of the Association for Computational Linguistics (ACL)*, 2021.
- [23] Timo Schick, Jane Dwivedi-Yu, Roberto Dessì, Roberta Raileanu, Maria Lomeli, Eric Hambro, Luke Zettlemoyer, Nicola Cancedda, and Thomas Scialom. Toolformer: Language models can teach themselves to use tools. In *Advances in Neural Information Processing Systems (NeurIPS)*, 2023.
- [24] Richard T. Snodgrass. *Developing Time-Oriented Database Applications in SQL*. Morgan Kaufmann, 1999.
- [25] Huanhuan Wu, James Cheng, Silu Huang, Yiping Ke, Yi Lu, and Yanyan Xu. Path problems in temporal graphs. *Proceedings of the VLDB Endowment*, 7(9):721–732, 2014.
- [26] Shunyu Yao, Jeffrey Zhao, Dian Yu, Nan Du, Izhak Shafran, Karthik Narasimhan, and Yuan Cao. ReAct: Synergizing reasoning and acting in language models. In *International Conference on Learning Representations (ICLR)*, 2023.
- [27] Tao Yu, Rui Zhang, Kai Yang, Michihiro Yasunaga, Dongxu Wang, Zifan Li, James Ma, Irene Li, Qingning Yao, Shanelle Roman, Zilin Zhang, and Dragomir Radev. Spider: A large-scale human-labeled dataset for complex and cross-domain semantic parsing and text-to-SQL task. In *Proceedings of the 2018 Conference on Empirical Methods in Natural Language Processing (EMNLP)*, 2018.